

Capturing the Ebb and Flow: A Dynamic Momentum Model for Tennis Match Analysis

Summary

Tennis momentum swings are frequently decisive yet remain hard to quantify. This paper develops a comprehensive **Dynamic Momentum Score (DMS)** framework to identify, validate, and predict match flow. Utilizing point-level data from the **2023 Wimbledon Championships (7,284 points)**, we established a recursive scoring system based on **ELO-rating principles**. The model integrates a serve-advantage calibration (weighted by the observed **67.31%** win rate), key-point multipliers, and a time-decay mechanism. **Results indicate that the average DMS achieves a 90.3% accuracy in predicting match winners.**

To verify momentum as a physical reality, we conducted a **statistical validation suite**, including **Runs Tests** and **Conditional Probability Analysis**. The data-driven analysis reveals that a player's winning probability after a successful point rises to **54.4%**, significantly exceeding the **51.0%** baseline win rate ($p < 0.001$). These findings provide robust evidence that momentum is a **statistically significant, non-random phenomenon** in elite tennis rather than a simple random walk.

For predicting momentum "swings," we developed a **Random Forest (RF) Classifier** utilizing a multi-dimensional feature matrix of DMS gradients and rally counts. The predictive framework achieved a high **Area Under the Curve (AUC) of 0.945**. Feature importance identification confirms that "current momentum state" and "rally length" are the strongest predictors of impending shifts. **Leave-one-match-out cross-validation further confirms robust generalization with an AUC of 0.941.**

Finally, **sensitivity analysis** on decay constants and weights ensures model stability across diverse conditions. Our study concludes that while momentum is persistent, it is highly susceptible to "pivotal points." We advise that psychological intervention is most critical during high-rally-count points where DMS gradients fluctuate. **The proposed framework is versatile and can be generalized to other point-based sports to optimize strategic decision-making.**

Keywords: Tennis Momentum; Dynamic Scoring Model; Statistical Validation; Random Forest; Leave-One-Out Cross-Validation

Contents

1	Introduction	4
1.1	Problem Background	4
1.2	Restatement of Problems	4
1.3	Our Work	4
2	Preparation for Modeling	5
2.1	Model Assumptions	5
2.2	Notations	6
2.3	Data Preprocessing	6
3	Problem 1: Momentum Quantification Model	7
3.1	Model Overview	7
3.2	Model Formulation	7
3.2.1	Serve Advantage Factor	7
3.2.2	Key Point Multiplier	8
3.2.3	Streak Bonus	8
3.2.4	Decay Mechanism	8
3.3	Results and Visualization	9
3.4	Model Validation	9
4	Problem 2: Statistical Validation of Momentum	10
4.1	Runs Test for Randomness	10
4.2	Conditional Probability Analysis	11
4.3	Chi-Square Test	12
4.4	Conclusion on Momentum Reality	12
5	Problem 3: Momentum Shift Prediction	12
5.1	Problem Definition	12
5.2	Feature Selection	13

5.3	Model Selection and Comparison	14
5.4	Model Performance	15
5.5	Feature Importance Analysis	15
5.6	Practical Recommendations for Coaches	16
6	Problem 4: Model Generalization	17
6.1	Leave-One-Match-Out Cross-Validation	17
6.2	Performance Across Tournament Rounds	17
6.3	Applicability to Other Contexts	18
7	Sensitivity Analysis	19
7.1	Momentum Model Parameters	19
7.2	Random Forest Hyperparameters	20
7.3	Robustness Conclusion	21
8	Model Analysis	21
8.1	Strengths	21
8.2	Weaknesses and Future Work	21
9	Memorandum	23
10	References	24

1 Introduction

1.1 Problem Background

In the 2023 Wimbledon Men's Singles Final, 20-year-old Carlos Alcaraz defeated 36-year-old Novak Djokovic in a match characterized by remarkable momentum swings. Djokovic dominated the first set 6-1, yet Alcaraz captured the second set 7-6 in a tiebreak. The third set mirrored the first in reverse, with Alcaraz winning 6-1. Djokovic then rallied to take the fourth set 6-3 before Alcaraz ultimately prevailed 6-4 in the fifth set.

Such dramatic fluctuations in performance, where one player appears to gain "momentum" or "force" over their opponent, are frequently observed in competitive sports. While commentators and coaches often reference momentum as a decisive factor, quantifying this phenomenon and understanding its underlying drivers remain significant challenges.

1.2 Restatement of Problems

The problem requires us to address four key questions:

1. **Momentum Quantification:** Develop a model to capture and visualize match flow, accounting for serve advantage.
2. **Momentum Validation:** Test whether momentum is a real phenomenon or merely random fluctuation.
3. **Momentum Prediction:** Build a predictive model for momentum shifts and identify key influencing factors.
4. **Model Generalization:** Evaluate model performance across different matches and discuss applicability to other sports.

1.3 Our Work

To address these challenges, we develop a comprehensive analytical framework:

- We construct a **Dynamic Momentum Score (DMS)** model that updates after each point, incorporating serve advantage, key point multipliers, streak bonuses, and decay mechanisms.
- We employ **statistical hypothesis testing** (runs test, conditional probability, chi-square test) to validate momentum as a non-random phenomenon.
- We train a **Random Forest classifier** to predict momentum shifts, using feature importance analysis to identify key factors.
- We conduct **leave-one-match-out cross-validation** to assess generalization capability and discuss broader applicability.




Figure 1: Our Work

2 Preparation for Modeling

2.1 Model Assumptions

Assumption 1: The provided data accurately reflects match events.

→ **Justification:** The dataset is officially compiled from Wimbledon Championships records with detailed point-by-point tracking via Hawk-Eye technology.

Assumption 2: Serve advantage remains relatively constant throughout a match.

→ **Justification:** Analysis of our dataset shows servers win 67.31% of points (4,903/7,284). We verified stability by computing per-set serve win rates, which ranged from 64.8% to 69.2% across all sets, confirming the assumption. We use $p_{\text{serve}} = 0.65$ as a conservative approximation within this range.

Assumption 3: External factors (weather, crowd influence) affect both players equally.

→ **Justification:** Wimbledon's controlled environment and retractable roof minimize external variations. This assumption is admittedly strong for outdoor court matches.

Assumption 4: Player fatigue affects both competitors similarly over the match duration.

→ **Justification:** Professional players undergo comparable physical conditioning. How-

ever, this assumption may not hold for significantly mismatched opponents or in fifth-set tiebreaks.

Limitation: Assumptions 3–4 are not directly validated in this study. Future work could incorporate player-specific fatigue indicators and external condition data.

2.2 Notations

Table 1: Notation definitions with parameter justifications.

Symbol	Description	Justification
M_t	Momentum score after point t	State variable
ΔM_t	Momentum change at point t	State variable
w_{base}	Base weight (1.0)	Normalization constant
p_{serve}	Serve advantage (0.65)	Data: actual serve win rate = 67.31%; grid search: 0.65–0.75 optimal
w_{break}	Break point multiplier (1.5)	Grid search: 1.5–2.0 achieves max accuracy
w_{key}	Key point multiplier (1.2)	Domain knowledge: set/match points are critical but less frequent than break points; value chosen below w_{break}
γ	Decay rate (0.02)	Sensitivity analysis: minimal impact; prevents infinite accumulation

2.3 Data Preprocessing

The dataset contains 7,284 points from 31 matches in the 2023 Wimbledon Men’s Singles (Round 2 onwards). We performed the following preprocessing steps:

1. **Missing Value Treatment:** Table 2 summarizes missing data statistics. For `return_depth` (1,309 missing, 17.97%), values were filled with “ND” (mode) when not an ace. The high missing rate is due to unreturned serves.

Table 2: Missing value summary.

Variable	Missing	Rate	Imputation
<code>return_depth</code>	1,309	17.97%	Mode (“ND”)
<code>speed_mph</code>	752	10.32%	Not used in model
<code>serve_width</code>	54	0.74%	Mode
<code>serve_depth</code>	54	0.74%	Mode

2. **Feature Engineering** We created derived features as shown in Table 3.

Table 3: Derived Features for Tennis Match Analysis

Feature Name	Description
global_point_idx	Sequential point number within each match
point_diff	Cumulative point difference between players
p1_rolling_win_rate_5/10	Rolling win rates with 5 and 10 point windows
p1_streak/p2_streak	Current winning streak length
is_break_point/is_key_point	Binary indicators for critical points
point_duration	Time elapsed since previous point

Note: We tested window sizes of 3, 5, 7, 10, 15, and 20 points for the `p1_rolling_win_rate_5/10` feature. Correlation analysis between rolling win rate and next-point outcome shows that window size 7 achieves the strongest correlation ($r = -0.128$), followed by window 5 ($r = -0.101$). We chose window 5 as the primary feature because it (a) has strong predictive correlation, (b) matches a typical game length in tennis (46 points), and (c) provides sufficient responsiveness to recent performance changes.

3. **Data Type Conversion:** Elapsed time converted to seconds; categorical variables converted to category type.

3 Problem 1: Momentum Quantification Model

3.1 Model Overview

To capture the “ebb and flow” of tennis matches, we develop a Dynamic Momentum Score (DMS) model inspired by the ELO rating system. The model quantifies momentum as a continuous variable updated after each point, with positive values indicating Player 1’s advantage and negative values indicating Player 2’s advantage.

3.2 Model Formulation

The momentum score is updated according to:

$$M_t = M_{t-1} + \Delta M_t - \gamma \cdot M_{t-1} \quad (1)$$

where $M_0 = 0$ (balanced start), and the momentum change ΔM_t is computed as:

$$\Delta M_t = w_{\text{base}} \cdot w_{\text{serve}} \cdot w_{\text{key}} \cdot w_{\text{streak}} \cdot \text{sign}(v_t) \quad (2)$$

Here, $v_t = 1$ if Player 1 wins point t , and $v_t = -1$ otherwise.

3.2.1 Serve Advantage Factor

The serve advantage factor adjusts momentum changes based on whether the server won:

$$w_{\text{serve}} = \begin{cases} p_{\text{serve}}, & \text{if server wins} \\ 1/p_{\text{serve}}, & \text{if receiver wins} \end{cases} \quad (3)$$

Mathematical Derivation: The reciprocal form arises from probability odds. If the server wins with probability p , then:

- Expected server win odds: $\frac{p}{1-p}$
- A server win is “expected,” so we weight it by p (below 1.0)
- A receiver win has odds $\frac{1-p}{p}$; we approximate this “surprise factor” as $1/p$

This ensures that breaking serve (receiver winning) contributes approximately $1/0.65 \approx 1.54$ times the momentum of holding serve, reflecting the psychological significance of break points. Grid search confirms $p_{\text{serve}} \in [0.65, 0.75]$ achieves maximum accuracy (90.3%).

3.2.2 Key Point Multiplier

Critical points receive higher weight:

$$w_{\text{key}} = \begin{cases} w_{\text{break}}, & \text{if break point} \\ 1.2, & \text{if set point or match point} \\ 1.0, & \text{otherwise} \end{cases} \quad (4)$$

3.2.3 Streak Bonus

Consecutive wins amplify momentum changes:

$$w_{\text{streak}} = 1 + \beta \cdot \text{streak_length} \quad (5)$$

We set $\beta = 0.1$ based on sensitivity analysis, which shows prediction accuracy remains stable across $\beta \in [0.0, 0.2]$ (accuracy varies $< 1\%$). The value 0.1 provides moderate amplification: a 3-point streak increases weight by 30%, capturing the “hot hand” effect without over-weighting.

3.2.4 Decay Mechanism

The term $-\gamma \cdot M_{t-1}$ ensures momentum naturally decays toward zero, preventing indefinite accumulation.

Why Linear Decay: We chose linear (first-order) decay over exponential decay for two reasons:

1. **Simplicity:** Linear decay has a single interpretable parameter.
2. **Robustness:** Sensitivity analysis shows that decay rate has minimal impact on accuracy (all tested values $\gamma \in [0.0, 0.05]$ achieve 90.3% accuracy), so the simpler form suffices.

The default $\gamma = 0.02$ means momentum decays by 2% per point, roughly halving over 35 points (approximately one set).

3.3 Results and Visualization

We applied the DMS model to all 31 matches. Figure 2 displays the momentum trajectory for the 2023 Wimbledon Final between Alcaraz and Djokovic.




Figure 2: Momentum curve for the 2023 Wimbledon Final

The visualization reveals clear momentum swings corresponding to each set's outcome. Djokovic's dominance in Set 1 (reaching momentum ≈ -12) is followed by Alcaraz's recovery in Sets 2-3 (peak momentum $\approx +25$).




Figure 3: Momentum heatmap

Figure 3 provides a complementary view of momentum distribution across each set, highlighting the intensity and duration of momentum swings.

3.4 Model Validation

To assess predictive validity, we computed average momentum for each match and compared it with actual winners:

Table 4: Momentum model performance across all matches

Metric	Value
Matches Analyzed	31
Total Points	7,284
Correct Predictions	28
Prediction Accuracy	90.3%




Figure 4: Relationship between average momentum and match outcome

Matches with positive average momentum predominantly resulted in Player 1 victories, confirming the model's discriminative power. As illustrated in Figure 4

4 Problem 2: Statistical Validation of Momentum

A skeptical coach argues that momentum might simply be random fluctuation. We employ three statistical tests to evaluate this hypothesis.

4.1 Runs Test for Randomness

The runs test examines whether the sequence of point winners exhibits more clustering than expected under randomness.

Let R denote the observed number of runs (consecutive sequences of wins by the same player), with expected value under randomness:

$$E(R) = \frac{2n_1n_2}{n_1 + n_2} + 1 \quad (6)$$

where n_1 and n_2 are the total points won by each player.




Figure 5: Distribution of runs ratio (observed/expected) across all matches

Results show that the average runs ratio is **0.935**, significantly below 1.0 ($p < 0.05$), indicating that winning streaks cluster more than random chance would predict. Figure 5 visualizes the distribution of runs ratios across all matches.

4.2 Conditional Probability Analysis

We compare the probability of winning after a previous win versus the overall winning probability:

$$P(\text{Win}_t | \text{Win}_{t-1}) \quad \text{vs} \quad P(\text{Win}) \quad (7)$$

Calculation Scope: We analyze all 7,284 points across 31 matches, using lagged outcomes to avoid within-point correlation. The analysis distinguishes serve/receive contexts:

Table 5: Conditional probability analysis by serve context.

Condition	P1 Win Rate	Sample Size
Overall (P1)	51.07%	7,253
P1 Won Last Point	54.35%	3,675
P1 Lost Last Point	47.65%	3,578
P1 Serving	68.72%	—
P1 Receiving	34.10%	—

Key findings (visualized in Figure 6):

- $P(\text{Win} | \text{Won Last}) = 54.35\%$
- $P(\text{Win}) = 51.07\%$
- Difference: **3.29 percentage points**
- Server advantage effect: P1 wins 68.72% when serving vs 34.10% when receiving, confirming the importance of controlling for serve in momentum analysis.




Figure 6: Conditional probability comparison

4.3 Chi-Square Test

A chi-square test confirms the statistical significance of the conditional probability difference:

$$\chi^2 = \sum \frac{(O_i - E_i)^2}{E_i} \quad (8)$$

The test yields $p < 0.001$, strongly rejecting the null hypothesis that previous point outcomes do not influence subsequent performance.

4.4 Conclusion on Momentum Reality

All three statistical tests provide consistent evidence that momentum in tennis is a **real phenomenon**, not merely random fluctuation. The “hot hand effect” is statistically significant, with players more likely to continue winning after recent success.

5 Problem 3: Momentum Shift Prediction

5.1 Problem Definition

We define a “momentum shift” as a sign change in the momentum score, occurring when:

$$M_t \cdot M_{t-1} < 0 \quad \text{and} \quad |M_{t-1}| > \tau \quad (9)$$

where $\tau = 1.0$ is a threshold to filter minor fluctuations.

Threshold Selection: We tested $\tau \in \{0.5, 0.75, 1.0, 1.25, 1.5, 2.0\}$ via sensitivity analysis. Table 6 summarizes the results.

Table 6: Threshold sensitivity analysis for momentum shift definition.

Threshold τ	Shift Events	Shift Rate	CV AUC
0.50	271	3.74%	0.920
0.75	150	2.07%	0.916
1.00	106	1.46%	0.915
1.25	70	0.97%	0.888
1.50	31	0.43%	0.843
2.00	5	0.07%	N/A

Figure 7 visualizes the trade-off between threshold stringency and model performance. The left panel shows how the number of detected shift events decreases as the threshold increases, reflecting a more restrictive definition of meaningful momentum shifts. The right panel demonstrates that prediction AUC remains relatively stable across a range of thresholds but degrades when the sample size becomes insufficient at high thresholds.




Figure 7: Threshold sensitivity analysis for momentum shift definition. Left: Number of detected shift events vs threshold value. Right: Cross-validation AUC vs threshold value, showing stable performance up to $\tau = 1.25$.

Key observations:

- $\tau = 0.5$: High AUC (0.920) but includes minor fluctuations that may not represent meaningful momentum shifts.
- $\tau = 1.0$: Optimal balance—sufficient samples (106 events) with competitive AUC (0.915). The threshold filters out noise while retaining psychologically significant shifts.
- $\tau \geq 1.5$: Insufficient samples for reliable training; AUC degrades.

The choice $\tau = 1.0$ balances signal quality (filtering noise) with sample size requirements for robust machine learning.

5.2 Feature Selection

Based on domain knowledge and data availability, we select 14 predictive features:

Table 7: Features used for momentum shift prediction.

Category	Features
Match Progress	set_no, games_in_set, sets_played
Score State	point_diff, momentum_prev
Player Performance	p1_streak_prev, p2_streak_prev, p1_rolling_win_rate_5
Serve Context	is_p1_serving, serve_no
Point Criticality	is_break_point, is_key_point
Rally Characteristics	rally_count, point_duration

5.3 Model Selection and Comparison

We compared three classifiers using 5-fold cross-validation:

Table 8: Model comparison (5-fold CV AUC).

Model	Mean AUC	Std AUC
Logistic Regression	0.627	0.026
Random Forest	0.915	0.026
Gradient Boosting	0.987	0.004

While Gradient Boosting achieves the highest AUC (0.987), we selected **Random Forest** for final deployment because:

1. **Interpretability:** Feature importance is directly available via Gini impurity, enabling actionable coaching insights. Gradient Boosting feature importance is less intuitive.
2. **Overfitting Concern:** The extremely high AUC (0.987) with very low variance (0.004) suggests potential overfitting, especially given our small dataset (31 matches, ~7,200 points). Random Forest's more moderate AUC with comparable variance indicates better generalization potential.
3. **Training Efficiency:** Parallel tree construction allows faster experimentation during hyperparameter tuning.

The test-set AUC of 0.945 reported below confirms Random Forest's strong generalization despite the lower cross-validation AUC.

Final Random Forest configuration:

- Number of trees: 100 (stabilized AUC beyond 50 trees)
- Maximum depth: 10 (grid search: depth >10 shows diminishing returns)
- Minimum samples for split: 20
- Class weight: balanced (to handle class imbalance)

5.4 Model Performance




Figure 8: ROC curve for momentum shift prediction

The model achieves strong predictive performance (see Figure 8):

- **Test AUC: 0.945**
- **5-Fold CV AUC: 0.924 ± 0.022**

5.5 Feature Importance Analysis

Figure 9 and Table 9 present the complete feature importance ranking from the Random Forest model.




Figure 9: Feature importance ranking from Random Forest

Table 9: Complete feature importance ranking.

Rank	Feature	Importance
1	momentum_prev	0.443
2	rally_count	0.145
3	point_diff	0.116
4	point_duration	0.064
5	games_in_set	0.044
6	sets_played	0.035
7	set_no	0.033
8	p1_rolling_win_rate_5	0.026
9	p1_streak_prev	0.024
10	p2_streak_prev	0.022
11	serve_no	0.018
12	is_p1_serving	0.016
13	is_key_point	0.008
14	is_break_point	0.006

Key insights:

1. **Previous Momentum** (0.443): Dominates prediction; extreme momentum values are prone to mean reversion
2. **Rally Count** (0.145): Longer rallies create more uncertainty and shift opportunities
3. **Point Difference** (0.116): Large score gaps may indicate unstable equilibrium
4. **Point Duration** (0.064): Longer points correlate with competitive exchanges
5. **Surprisingly**: Break/key point indicators rank low (0.006–0.008), suggesting their importance is already captured by other features

5.6 Practical Recommendations for Coaches

Based on our analysis, we provide the following strategic recommendations:

1. **Prioritize Break Points**: Break points show $1.2\times$ higher probability of momentum shifts. Prepare players mentally for these critical moments.
2. **Maintain Focus During Winning Streaks**: Long winning streaks become increasingly vulnerable. Players should avoid complacency.
3. **Stay Resilient When Behind**: Momentum is inherently unstable in later sets. Falling behind does not preclude comeback.
4. **Disrupt Opponent Streaks Early**: Actively target opponent momentum by varying tactics during their winning runs.

6 Problem 4: Model Generalization

6.1 Leave-One-Match-Out Cross-Validation

To rigorously test generalization, we employ Leave-One-Match-Out (LOMO) cross-validation: for each of the 31 matches, we train on 30 matches and evaluate on the held-out match. Figure 10 displays the results for each match. Each bar represents AUC for a held-out match.




Figure 10: Leave-one-match-out cross-validation results

Table 10: LOMO cross-validation performance summary

Metric	Value
Mean AUC	0.941
Std AUC	0.032
Min AUC	0.856
Max AUC	0.988
Winner Prediction Accuracy	90.0%

The model demonstrates consistent performance across different match types and rounds, confirming robust generalization.

6.2 Performance Across Tournament Rounds

Table 11: Model performance by tournament round (LOMO cross-validation).

Round	Matches	Mean AUC	Win Pred. Acc.
Round 3	15	0.936	86.7%
Round 4 (R16)	8	0.954	87.5%
Quarterfinals	4	0.951	100%
Semifinals	2	0.963	100%
Final	1	0.988	100%

Note: Some matches in LOMO validation were excluded when they contained insufficient momentum shift events (<1 event) for testing.

Figure 11 provides a visual representation of model performance across tournament rounds. The bar chart displays individual match AUC scores, with green bars indicating correct winner prediction and red bars indicating incorrect predictions. The vertical dashed lines separate different tournament rounds, revealing a clear trend of improving performance in later stages.




Figure 11: Model performance by tournament round

Performance improves in later rounds. The final match (Alcaraz vs Djokovic) achieves the highest AUC (0.988), which we attribute to several factors quantified in our analysis:

1. **Larger sample size:** 334 points vs average of 232 points in other matches (+44%)
2. **Higher break point frequency:** 10.2% break point rate vs 6.8% average (+50%), providing stronger momentum signals
3. **Longer rallies:** Average rally count of 4.46 vs 3.06 in other matches (+46%), creating more dramatic momentum swings
4. **Greater momentum volatility:** Momentum standard deviation of 8.55 vs 5.95 average (+44%), making shifts more pronounced and easier to detect
5. **Five-set format:** Extended match duration allows momentum patterns to fully develop

6.3 Applicability to Other Contexts

We assessed the applicability of our momentum framework to other sports contexts based on structural similarity to tennis (point-by-point scoring, server advantage, etc.). Figure 12 and Table 12 summarize our assessment.




Figure 12: Estimated applicability scores for different sports and contexts.

Table 12: Model applicability assessment.

Context	Score	Notes
Other Men’s Tennis	High	Directly applicable with same parameters
Women’s Tennis	Medium-High	Requires serve advantage recalibration
Other Grand Slams	Medium-High	Surface differences may affect parameters
Table Tennis	Medium	Shorter rallies; model structure adaptable
Team Sports	Low	Requires multi-agent modeling

7 Sensitivity Analysis

7.1 Momentum Model Parameters

We assess the sensitivity of the DMS model to its four key parameters via grid search. Figure 13 visualizes the sensitivity landscape.

Table 13: Parameter grid search results.

Parameter	Tested Range	Optimal Range	Default	Accuracy
serve_advantage	0.55–0.75	0.65–0.75	0.65	87.1–90.3%
break_point_mult	1.0–2.5	1.5–2.0	1.5	87.1–90.3%
streak_bonus	0.0–0.2	0.0–0.2	0.1	90.3% (stable)
decay_rate	0.0–0.05	0.0–0.05	0.02	90.3% (stable)




Figure 13: Sensitivity of momentum model to parameter variations

Key findings:

- **Serve advantage:** Most sensitive parameter. Values below 0.60 reduce accuracy to 87.1%. Our data-derived value (0.6731) and default (0.65) both achieve optimal 90.3%.
- **Break point multiplier:** Moderate sensitivity. Values ≥ 1.5 achieve optimal accuracy.
- **Streak bonus & decay rate:** Minimal sensitivity. Accuracy remains 90.3% across all tested values, confirming these are secondary parameters.

7.2 Random Forest Hyperparameters

Cross-validation AUC remains stable across reasonable hyperparameter ranges:

- **Number of trees:** AUC stabilizes above 50 trees
- **Max depth:** Optimal around 10; deeper trees show marginal improvement
- **Min samples split:** Minimal impact on AUC




Figure 14: Summary of sensitivity indices for all model parameters.

7.3 Robustness Conclusion

Figure 14 summarizes the sensitivity indices across all model parameters. The sensitivity analysis confirms that our models are robust across parameter variations. Default parameter configurations lie within optimal performance regions, and prediction accuracy remains above 80% under reasonable perturbations.

8 Model Analysis

8.1 Strengths

- **Theoretical Foundation:** The DMS model is grounded in established rating theory (ELO) and incorporates domain-specific factors (serve advantage, key points).
- **High Predictive Accuracy:** The model achieves 90.3% accuracy in winner prediction and 0.945 AUC in shift prediction.
- **Interpretability:** Feature importance provides actionable insights for coaches.
- **Robust Generalization:** LOMO cross-validation confirms stable performance across unseen matches.
- **Statistical Validation:** Multiple hypothesis tests provide strong evidence for momentum reality.

8.2 Weaknesses and Future Work

- **Data Limitation:** Analysis is restricted to 2023 Wimbledon men's singles (31 matches, 7,284 points). Broader validation across tournaments and years would strengthen conclusions.
- **Player-Specific Effects:** The model does not account for individual player characteristics (playing style, mental resilience).

- **External Factors:** Weather conditions, crowd effects, and fatigue are not explicitly modeled.
- **Correlation vs. Causation:** Our statistical tests establish that momentum-like patterns exist (winning after winning is more likely than baseline). However, this does not prove that “feeling momentum” *causes* better performance. Alternative explanations include:
 - Fatigue asymmetry (one player tiring faster)
 - Tactical adjustments lagging behind play
 - Opponent demoralization (psychological effect on the *other* player)

Establishing true causality would require controlled experiments or instrumental variable analysis, which are beyond this paper’s scope.

Future work could incorporate player embeddings, real-time physiological data, and causal inference methods to address these limitations.

9 Memorandum

MEMORANDUM

To: Tennis Coaches and Athletic Directors

From: Mathematical Modeling Team

Subject: Understanding and Leveraging Momentum in Tennis Matches

Date: January 28, 2026

Executive Summary

Our analysis of 31 Wimbledon matches (7,284 points) reveals that momentum in tennis is a statistically validated phenomenon, not random noise. We have developed tools to quantify and predict momentum shifts, offering actionable insights for match preparation.

Key Findings

1. **Momentum is Real:** Statistical tests confirm that winning streaks cluster more than chance predicts. Players winning one point are 3.4% more likely to win the next.
2. **Break Points are Critical:** Break points show $1.2\times$ higher probability of triggering momentum shifts. Mental preparation for these moments is essential.
3. **Streaks are Fragile:** Extended winning streaks become increasingly vulnerable to interruption. Maintain focus even when dominating.
4. **Comebacks are Possible:** Later sets show greater momentum instability. Players should never concede mentally, regardless of score.

Recommendations

- Train players to recognize and manage momentum states during matches
- Develop specific routines for break points and other high-pressure situations
- Use tactical variation to disrupt opponent momentum during their winning runs
- Emphasize mental resilience training for late-set scenarios

Best Regards,
Team # 2613942

10 References

References

- [1] Braidwood, J. (2023). Novak Djokovic faced a unique opponent – does Wimbledon loss signal the beginning of the end? *The Independent*.
- [2] Merriam-Webster Dictionary. Momentum. <https://www.merriam-webster.com/dictionary/momentum>
- [3] Rivera, J. (2023). Tennis scoring explained: Wimbledon rules, terms, and scoring system. *Sporting News*.
- [4] Elo, A. E. (1978). *The Rating of Chessplayers, Past and Present*. Arco Publishing.
- [5] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5-32.
- [6] Gilovich, T., Vallone, R., & Tversky, A. (1985). The hot hand in basketball: On the misperception of random sequences. *Cognitive Psychology*, 17(3), 295-314.
- [7] Bar-Eli, M., Avugos, S., & Raab, M. (2006). Twenty years of “hot hand” research: Review and critique. *Psychology of Sport and Exercise*, 7(6), 525-553.

Report on Use of AI

AI Tools Used

We utilized AI assistance for the following tasks:

1. Claude (Anthropic) — Code Generation

- **Query:** “Write Python code to implement a momentum model for tennis with serve advantage weighting”
- **Output:** Initial class structure for `MomentumModel`
- **Modification Rate:** ~60%. We added break point multipliers, decay mechanism, and match-level evaluation functions.

2. Claude (Anthropic) — Statistical Analysis

- **Query:** “Implement runs test and chi-square test for sequence analysis”
- **Output:** Python functions for statistical tests
- **Modification Rate:** ~30%. Minor adjustments for our data format.

3. Cursor AI — Writing Assistance

- **Query:** LaTeX formatting, grammar checking, sentence restructuring

- **Output:** Draft text for methodology sections
- **Modification Rate:** ~70%. We rewrote most sections to add quantitative details, parameter justifications, and domain-specific interpretations.

Summary: AI tools accelerated initial code scaffolding and writing. All model design decisions, parameter tuning, statistical interpretation, and conclusions are original team work. Overall estimated human contribution: 75–80% of final content.